

Fusion-Core ISA Definition: Revision 0.1

Dylan Wadler

December 28, 2017

Contents

| | |
|---|----------|
| I Administrative | 4 |
| 1 Change log | 4 |
| 2 Introduction | 4 |
| 2.1 About | 4 |
| 2.2 Goals | 4 |
| 2.3 Conventions | 4 |
| II Programming Information | 5 |
| 3 Register File Definitions | 5 |
| 3.1 Register File List | 5 |
| 3.2 General Purpose Registers | 5 |
| 3.3 Special Registers | 6 |
| 3.4 Special Purpose Registers | 6 |
| 3.4.1 System Registers | 6 |
| 3.4.2 Supervisor Registers | 7 |
| 4 Permission Levels | 7 |
| 4.1 User Levels | 7 |
| 4.1.1 Low User Level | 7 |
| 4.1.2 High User Level | 8 |
| 4.2 Supervisor Levels | 8 |
| 4.2.1 Low Supervisor Level | 8 |
| 4.2.2 High Supervisor Level | 8 |
| 4.3 Hypervisor Levels | 8 |
| 4.3.1 Low Hypervisor Level | 8 |
| 4.3.2 High Hypervisor Level | 8 |
| 4.3.3 Top Hypervisor Level | 8 |
| 5 Memory | 8 |
| 5.1 Memory Locations for Vector Table | 8 |
| 5.1.1 Interrupt Vector Table | 8 |
| 5.1.2 Exception Vector Table | 8 |
| III Instructions | 9 |
| 6 Instruction Definitions | 9 |
| 6.1 Instruction Format Types | 9 |
| 6.1.1 Integer | 9 |
| 6.1.2 Immediate | 10 |
| 6.1.3 Load/Store | 10 |
| 6.1.4 Branch/Jump | 12 |
| 6.1.5 System Instructions | 13 |
| 6.1.6 Co-Processor | 13 |

| | |
|---|-----------|
| 6.2 List of OPCODEs | 14 |
| 6.3 Instruction Details | 14 |
| 6.3.1 Integer | 14 |
| 6.4 List of Instructions | 14 |
| 6.4.1 Immediate | 16 |
| 6.4.2 Load/Store | 16 |
| 6.4.3 Branch/Jump | 16 |
| 6.4.4 Co-Processor | 16 |
| 7 Exceptions and Interrupts | 16 |
| 7.1 Exceptions | 16 |
| 7.2 Interrupts | 16 |
| 7.2.1 User Level | 16 |
| 7.2.2 Supervisor Level | 16 |
| IV Co-Processors | 17 |
| 8 Co-Processor Overview | 17 |
| 8.1 Co-Processor Interface | 17 |
| 8.1.1 Co-Processor Conventions | 18 |
| 8.1.2 Register Connections | 18 |
| 8.1.3 Decode unit Connections | 18 |
| 8.2 Interface Connection Definitions | 18 |
| 8.3 Adding custom Co-Processor | 18 |
| 8.4 List of Co-Processors | 18 |
| 8.4.1 Floating Point | 18 |
| 8.4.2 System Unit | 18 |
| 8.4.3 Memory Management Unit | 18 |
| 8.4.4 Multiprocessor Communication Unit | 18 |
| 9 Global Register File | 18 |
| 10 Recommended Co-Processors | 18 |
| 10.1 Floating Point Unit | 19 |
| 10.1.1 Registers | 19 |
| 10.1.2 Instructions | 19 |
| 10.2 System Unit | 19 |
| 10.2.1 Registers | 19 |
| 10.2.2 Instructions | 19 |
| 10.3 Memory Management Unit | 19 |
| 10.3.1 Registers | 19 |
| 10.3.2 Instructions | 19 |
| 10.4 Multiprocessor Communication Unit | 19 |
| 10.4.1 Registers | 19 |
| 10.4.2 Instructions | 19 |

Part I

Administrative

1 Change log

Version 0.1 Initial Definition of the Instruction Set Architecture

2 Introduction

2.1 About

Introduction The Fusion-Core ISA is dedicated to creating an easily expandible architecture without having to recompile a program's binary. By use of defining an easy interface with a simple core instruction set, this allows for more freedom in implementation. High end processors and microcontrollers would only require slight variations in configuration, as their core would remain identical save for easy to maintain and scalable co-processors.

Main Ideas The architecture is Big endian, with a core instruction set that is RISC, but the co-processors do not need to adhere to the RISC philosophy. This allows for more flexibility in design, and possibly faster core clock speeds as the pipeline would depend on smaller amounts of logic. Only the instructions provided in this document are to be implemented in the main processor. The co-processors defined in this document are recommended, but not required for normal function. Co-processor documentation is to be provided by the creator, and should adhere to the standards of clarity and conciseness such that it can be easily implemented from the documentation alone in a HDL.

64 Bit instructions: At this moment in time, the Fusion-Core ISA is only a 32 bit ISA. Due to the focus on co-processors, older implementations could easily be modified to include 64 bit operations.

Co-Processors The Co-Processor interface is currently defined by setting the MSB within the OP Code field of an instruction, to decrease complexity of the Decode unit. In doing so, this allows for coprocessor code to be written in the same memory space as the main processor code. In the current iteration, up to 32 different coprocessors can be used, with the option for dynamic or static allocation of the OP Codes. The interface for the coprocessors is explained further in the dedicated section.

2.2 Goals

The main goal is to provide an architecture with a simple decoding unit and the ability to utilize a single binary for all implementations of the architecture.

2.3 Conventions

Document Conventions: Example code will be shown with monospace text. General purpose registers will be denoted with \$R# where # is the number of the register. Special purpose registers will be written with **bold** text.

Part II

Programming Information

3 Register File Definitions

This section goes over the different registers available in the ISA. Each register file name begins with "REGF", such as the first General Purpose Register File being REGFGP0. Any additional register files require the number after the name of the register file. Register files with additional numbers after them are bank switched to reduce space, hence why the number is required to denote the register file space used.

To alleviate context switching delays, three general purpose register files are bank available by bank switching. Only these register files are required for a minimal system, though more can certainly be implemented if required. Each register file can be accessed by changing the value in the BSELRF register.

3.1 Register File List

Figure 1: General Purpose Registers

| REGFGP0 | | REGFSYSC0 | |
|----------|---------------|-----------|---------------|
| Register | Register Name | Register | Register Name |
| \$R0 | ZER0 | \$R0 | ZER0 |
| \$R1 | SP0 | \$R1 | SP1 |
| \$R2 | FP0 | \$R2 | FP1 |
| \$R3 | GP0 | \$R3 | GP1 |
| \$R4 | RA0 | \$R4 | RA1 |
| \$R5 | ARG0 | \$R5 | SYSARG0 |
| \$R6 | ARG1 | \$R6 | SYSARG1 |
| \$R7 | ARG2 | \$R7 | SYSARG2 |
| \$R8 | ARG3 | \$R8 | SYSARG3 |
| \$R9 | RVAL0 | \$R9 | SYSARG4 |
| \$R10 | RVAL1 | \$R10 | SYSARG5 |
| \$R11 | GR0 | \$R11 | SYSRVAL0 |
| \$R12 | GR1 | \$R12 | SYSRVAL1 |
| \$R13 | GR2 | \$R13 | SYSRVAL2 |
| \$R14 | GR3 | \$R14 | SYSRVAL3 |
| \$R15 | GR4 | \$R15 | SYSRVAL4 |
| \$R16 | GR5 | \$R16 | GPR0 |
| \$R17 | GR6 | \$R17 | GPR1 |
| \$R18 | GR7 | \$R18 | GPR2 |
| \$R19 | GR8 | \$R19 | GPR3 |
| \$R20 | GR9 | \$R20 | GPR4 |
| \$R21 | GR10 | \$R21 | GPR5 |
| \$R22 | TMP0 | \$R22 | GPR6 |
| \$R23 | TMP1 | \$R23 | GPR7 |
| \$R24 | TMP2 | \$R24 | SYSTMPR0 |
| \$R25 | TMP3 | \$R25 | SYSTMPR1 |
| \$R26 | TMP4 | \$R26 | SYSTMPR2 |
| \$R27 | TMP5 | \$R27 | SYSTMPR3 |
| \$R28 | TMP6 | \$R28 | SYSTMPR4 |
| \$R29 | TMP7 | \$R29 | SYSTMPR5 |
| \$R30 | HI0 | \$R30 | SYSREGHI0 |
| \$R31 | LOW0 | \$R31 | SYSREGL0W0 |

3.2 General Purpose Registers

32 general purpose registers that are 32 bits wide are available, as shown in Figure 1, in the previous section. There is a distinction between the System Register File and the

Figure 2: Special Registers

| System Registers | | | |
|------------------|---------------------------|---------------|---------------|
| Register Name | Description | Width (bytes) | Address (Hex) |
| CPUREV | CPU Revision | 1 | 0x00000000 |
| CPNUM | Co-Processor Number | 1 | 0x00000001 |
| CP00 | Co-Processor 0 ID | 2 | 0x00000002 |
| CP01 | Co-Processor 1 ID | 2 | 0x00000004 |
| CP02 | Co-Processor 2 ID | 2 | 0x00000006 |
| CP03 | Co-Processor 3 ID | 2 | 0x00000008 |
| STAT | Status Register | 1 | 0x0000000a |
| n/a | RESERVED | 1 | 0x0000000b |
| OPCAR | Opcode Allocation Pointer | 4 | 0x0000000c |

Figure 3: Supervisor Registers

| Supervisor Registers | | |
|----------------------|---------------------------------|---------------|
| Register Name | Description | Address (Hex) |
| SMSTAT | Supervisor mode status register | 0x00000000 |
| PRCPR0 | Process Pointer register 0 | 0x00000004 |
| HTINFO | Hardware Thread Info | 0x00000008 |
| HTCTL | Hardware Thread Control | 0x0000000c |

General Purpose Register File, as during certain syscall instructions, the register files are bank switched. Only GP0 through GP7 are saved, for passing between the different banks.

While it is not defined by the architecture, larger general purpose registers can be used instead of 32 bit wide registers. The System Register File allows 8 bit addressing for the registers to be accessed, in order to utilize more of the memory space. If larger registers are needed, consider using a co-processor to for instructions that require larger operands. This provides code compatibility between different implementations.

3.3 Special Registers

The special registers are sorted between the System Registers and the Supervisor Registers. The system registers provide simple configuration values and some read-only registers to give the programmer information about the implementation. The registers defined in this manual are the bare minimum special purpose registers, and should be included for code compatibility.

The supervisor registers are aimed at higher level functions required for operating system environments. They are not essential for operation, and can either be partially implemented or not at all. The optional parts will be noted in the register descriptions.

3.4 Special Purpose Registers

3.4.1 System Registers

STAT: Status Register

| | | | | | | | |
|---|----|------|------|------|------|------|------|
| 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| Z | OV | PEMA | PML2 | PML1 | PML0 | INTN | SPCP |

| | | |
|------|--|-----------------------|
| STAT | Read only register for various processor state information. The flags are explained in detail below. | |
| Z | Zero Flag; Indicates whether the processed instruction's resulted in zero. Read Only The flag is set to 0 when the ALU calculation is 0, and 1 when the output is non-zero. | |
| OV | Overflow Flag; Indicates whether the processed instruction's result overflowed the 32 bit space. Read only. The overflow flag is set to the output of the carry out of the ALU. With addition, this would set the bit to a 1 when true. For subtraction, the inverse is true. | |
| PML | Permission Level; Indicates the running process' permission level. Read only. | |
| | PML | Description |
| | 000 | Low User Level |
| | 001 | High User Level |
| | 010 | Low Supervisor Level |
| | 011 | High Supervisor Level |
| | 100 | Low Hypervisor Level |
| | 101 | High Hypervisor Level |
| | 110 | Reserved |
| | 111 | Top Hypervisor Level |
| | The permission levels are explained in more detailed in their dedicated section. | |
| PEMA | Permission Accepted; Indicates whether a privileged systemcall request was accepted. Read only | |
| INTN | Interrupt Enable; Indicates whether interrupts are enabled. Read/Write | |
| SPCP | Support CoProcessors; Indicates whether co-processor code is recognized as illegal instruction or microcode operation. Read/Write. | |

In order to write to this register, the bits that are read only can be set to any value.

Optionality In the event that the permission levels are not needed, they should be hard coded to 0x7, the highest permission level to avoid porting code. PEMA should also be hardcoded to a logic high for the same reasons stated.

3.4.2 Supervisor Registers

4 Permission Levels

4.1 User Levels

4.1.1 Low User Level

User space programs are designed to run in this permission level. Write access to supervisor registers is revoked.

4.1.2 High User Level

4.2 Supervisor Levels

4.2.1 Low Supervisor Level

4.2.2 High Supervisor Level

4.3 Hypervisor Levels

4.3.1 Low Hypervisor Level

4.3.2 High Hypervisor Level

4.3.3 Top Hypervisor Level

5 Memory

At this point in time, the ISA only handles 32 bit addresses. With memory capacity increasing in size as time goes on, this may change.

5.1 Memory Locations for Vector Table

5.1.1 Interrupt Vector Table

5.1.2 Exception Vector Table

| Address (32 bit) | Definition |
|------------------|----------------------------------|
| 0x0000 | Reset |
| 0x0004 | Co-Processor Microcode Exception |
| 0x0008 | |
| 0x000c | |
| 0x0010 | |
| 0x0014 | |

Figure 4: Exception Vector Table

Part III

Instructions

6 Instruction Definitions

6.1 Instruction Format Types

This section will talk about the different instructions available in the core processor, their encodings, function, and hazards they cause or registers they affect in the processor. The Co-Processor instructions are purely generic and allow the implementation of each Co-Processor to determine how their respective instructions will be decoded. Only the Operation Codes will be defined for each Co-Processor slot to allow for more customization.

6.1.1 Integer

The integer instructions are the heart of this processor's arithmetic abilities and is vital to ensure fast execution. A semi-strict adherence to RISC philosophy in this architecture is required to exploit any benefits to this ISA in a real implementation.

The integer instruction coding with descriptions, is shown in the diagrams below.

Register/Integer Instruction Format

| opcode (0x01) | | rd | | rsa | | rsb | | shft | | aluop | |
|---------------|----|----|----|-----|----|-----|----|------|---|-------|---|
| 31 | 26 | 25 | 21 | 20 | 16 | 15 | 11 | 10 | 4 | 3 | 0 |

Overview: The Register/Integer Instruction Format is for basic ALU operations, without immediates. Registers RSa and RSb are the two operands, which are stored in register Rd. The 4 bit ALUOP field denotes the settings for the ALU, to reduce complexity of selecting what operation to choose. The shft bits are only for the shift amount with the shifting instructions, but unused for other instructions.

The following instructions use this encoding:

| Instruction | Description | ALU Operation (hex) |
|-------------|------------------------|---------------------|
| add | Add | 0x0 |
| sub | Subtract | 0x1 |
| addu | Add Unsigned | 0x2 |
| subu | Subtract Unsigned | 0x3 |
| tcmp | 2's Complement | 0x4 |
| and | And | 0x5 |
| or | Or | 0x6 |
| xor | Exclusive Or | 0x7 |
| sal | Arithmetic Shift Left | 0x8 |
| sar | Arithmetic Shift Right | 0x9 |
| sll | Logic Shift Left | 0xa |
| slr | Logic Shift Right | 0xb |

| | | |
|----------|---------|-----|
| comp | Compare | 0xc |
| Reserved | n/a | 0xd |
| Reserved | n/a | 0xe |
| Reserved | n/a | 0xf |

6.1.2 Immediate

Immediate Instruction Format

| opcode | | rd | | rsa | | Immediate | | aluop | |
|--------|----|----|----|-----|----|-----------|---|-------|---|
| 31 | 26 | 25 | 21 | 20 | 16 | 15 | 4 | 3 | 0 |

Overview: The Immediate Instruction Format is for ALU operations that require an immediate value. The immediate field is only 12 bits wide, so in the event that a larger value is required, the Load Immediate Format should be used. There is only a single source register, RSa, with the other source being the immediate value. The result is stored into register Rd.

The following instructions use this encoding:

| Instruction | Description | ALU Operation (hex) |
|-------------|------------------------------|---------------------|
| addi | Add Immediate | 0x0 |
| subi | Subtract Immediate | 0x1 |
| addui | Add Unsigned Immediate | 0x2 |
| subui | Subtract Immediate Immediate | 0x3 |
| andi | And Immediate | 0x5 |
| ori | Or Immediate | 0x6 |
| xori | Exclusive Or Immediate | 0x7 |
| sali | Arithmetic Shift Left | 0x8 |
| sari | Arithmetic Shift Right | 0x9 |
| slli | Logic Shift Left | 0xa |
| slri | Logic Shift Right | 0xb |
| compi | Compare Immedaite | 0xc |

6.1.3 Load/Store

Load Instruction Format

| opcode | | rd | | rsa | | funct | | Immediate | |
|--------|----|----|----|-----|----|-------|----|-----------|---|
| 31 | 26 | 25 | 21 | 20 | 16 | 15 | 14 | 13 | 0 |

Overview: The Load Instruction Format is for reading values from memory into a register. A 14 bit immediate is used for a relative address calculation with RSa as the base address. The value is then stored into register Rd. The funct field is used to determine the byte size to read from memory. This encoding is described in the table below.

| Funct | Description |
|-------|----------------------------|
| 00 | 1 word (32 bits) |
| 01 | half word (16 bits, upper) |
| 10 | 24 bits, upper |
| 11 | byte, upper |

The following instructions use this encoding:

| Instruction | Description | Funct |
|-------------|-----------------------------|-------|
| lw | Load Word | 0x0 |
| lh | Load Half Word | 0x1 |
| lb | Load Byte | 0x3 |
| lth | Load Three (Bytes, 24 bits) | 0x2 |

Load Immediate Instruction Format

| opcode | rd | DSEL | Immediate |
|--------|-------|-------|-----------|
| 31 26 | 25 21 | 20 16 | 15 0 |

Overview: The Load Immediate Instruction Format is for loading immediate values into a register. The DSEL value determines the destination and function of the instruction. This is to allow the selection of various register files to be updated. The DSEL bit values and their function are shown below. The MSB of DSEL determines if the value is unsigned (1) or signed (0).

| DSEL (hex) | Destination |
|------------|--|
| 0x0000 | General purpose register file |
| 0x0001 | System Register File |
| 0x0010 | Global Co-Processor Register File |
| 0x0X1X | Reserved |
| 0x1000 | General purpose register file (Unsigned) |
| 0x1001 | System Register File (Unsigned) |
| 0x1010 | Global Co-Processor Register File (Unsigned) |

The following instructions use this encoding:

| Instruction | Description | DSEL |
|-------------|-----------------------|--------|
| li | load immediate (gprf) | 0x0000 |
| lsi | load immediate system | 0x0001 |
| lgi | load immediate global | 0x0002 |
| lui | load immediate (gprf) | 0x0008 |
| lsui | load immediate system | 0x0009 |
| lgui | load immediate global | 0x000a |

Store Instruction Format

| opcode | | Funct | | Immediate[13:11] | | rsa | rsb | Immediate[10:0] | | | |
|--------|----|-------|----|------------------|----|-----|-----|-----------------|----|----|---|
| 31 | 26 | 25 | 24 | 23 | 21 | 20 | 16 | 15 | 11 | 10 | 0 |

Overview: The Store Instruction Format is for writing register values to memory. A 14 bit immediate is used for a relative address calculation with RSa as the base address. The value in RSb is then written to memory. The funct field is used to determine the byte size to write to memory. This encoding is described in the table below.

| Funct | Description |
|-------|----------------------------|
| 00 | 1 word (32 bits) |
| 01 | half word (16 bits, upper) |
| 10 | 24 bits, upper |
| 11 | byte, upper |

The following instructions use this encoding:

| Instruction | Description | Funct |
|-------------|------------------------------|-------|
| sw | Store Word | 0x0 |
| sh | Store Half Word | 0x1 |
| sb | Store Byte | 0x3 |
| sth | Store Three (Bytes, 24 bits) | 0x2 |

6.1.4 Branch/Jump

Jump Instruction Format

| opcode | | Immediate[20:16] | | rsa | Immediate[15:0] | | |
|--------|----|------------------|----|-----|-----------------|----|---|
| 31 | 26 | 25 | 21 | 20 | 16 | 15 | 0 |

Overview: The Jump Instruction format is for changing the Program Counter. All Jumps are relative, except Jump Register Instructions, due to the non byte aligned size (20 bits) of the immediate field. Jumps use the value in register RSa to calculate the new PC. For pure relative jumps, the Zero register is used as the base. For jump instructions that link the program counter, the PC is stored in RA0, or R4 in the General Purpose Register file.

The following instructions use this encoding:

| Instruction | Description |
|-------------|------------------------|
| j | Jump |
| jal | Jump and Link |
| jr | Jump Register |
| jrl | Jump Register and Link |

Branch Instruction Format

| opcode | | Immediate[13:9] | | rsa | rsb | Immediate[8:0] | | funct | |
|--------|----|-----------------|----|-----|-----|----------------|----|-------|---|
| 31 | 26 | 25 | 21 | 20 | 16 | 15 | 11 | 10 | 0 |

Overview: The Branch Instruction Format is used for conditionally changing the Program Counter. The PC is only updated with whichever operation is indicated as being true or false.

The following instructions use this encoding:

| Instruction | Description | Funct |
|-------------|------------------------|-------|
| beq | Branch if Equal | 00 |
| bne | Branch if not Equal | 01 |
| bgt | Branch if Greater than | 10 |
| blt | Branch if Less than | 11 |

6.1.5 System Instructions

System Instruction Format

| opcode | rd | rsa | Funct | Immediate[7:0] |
|--------|-------|-------|-------|----------------|
| 31 26 | 25 21 | 20 16 | 15 8 | 7 0 |

Overview: The System Instruction Format is used for various system controls. Their functions vary, and are described in more detail in their individual listings.

The following instructions use this encoding:

| Instruction | Description | Funct | Uses Immediate |
|-------------|-----------------------------------|-------|----------------|
| syscall | System Call | 0x00 | yes |
| sysret | System Return | 0x01 | no |
| stspr | Store Special Purpose Register | 0x02 | yes |
| ldspr | Load Special Purpose Register | 0x03 | yes |
| sync | Synchronize memory/Flush Pipeline | 0x04 | no |
| lock | Lock Memory | 0x05 | yes |
| test | Test Lock | 0x06 | yes |
| pmir | Permission Increase Request | 0x07 | yes |
| pmd | Permission Decrease | 0x08 | yes |

6.1.6 Co-Processor

Since the coprocessors can have any implementation, only the opcode is required. However, adhering to similar formatting of the instruction formats provided in previous sections is imperative and the Fusion-Core foundation will not provide an accepted coprocessor ID number. More information about coprocessors can be accessed in the Co-Processor section.

6.2 List of OPCODEs

| Op Code | Description | Instrucion Formats |
|---------|-------------|--------------------|
|---------|-------------|--------------------|

6.3 Instruction Details

This section will go into detail about the instructions that are available, with their functions and options.

6.3.1 Integer

6.4 List of Instructions

Instruction Instruction Summary Table

| Integer Instructions | | |
|------------------------|---|---------------------------------|
| Instruction | Function | Binary |
| add | $Rd = RSa + RSb$ | 000001dddddaaaaabbbbxxxxxx0000 |
| sub | $Rd = RSa - RSb$ | 000001dddddaaaaabbbbxxxxxx0001 |
| addu | $Rd = RSa + RSb$ (Unsigned) | 000001dddddaaaaabbbbxxxxxx0010 |
| subu | $Rd = RSa - RSb$ (Unsigned) | 000001dddddaaaaabbbbxxxxxx0011 |
| tcmp | $Rd = !Rs_a$ | 000001dddddaaaaaxxxxxxxxxxx0100 |
| and | $Rd = RSa \& RSb$ | 000001dddddaaaaabbbbxxxxxx0101 |
| or | $Rd = RSa RSb$ | 000001dddddaaaaabbbbxxxxxx0110 |
| xor | $Rd = RSa \oplus RSb$ | 000001dddddaaaaabbbbxxxxxx0111 |
| sal | $Rd = RSa \ll RSb$ | 000001dddddaaaaabbbbssssss1000 |
| sar | $Rd = RSa \gg RSb$ | 000001dddddaaaaabbbbssssss1001 |
| sll | $Rd = RSa \ll RSb$ | 000001dddddaaaaabbbbssssss1010 |
| slr | $Rd = RSa \gg RSb$ | 000001dddddaaaaabbbbssssss1011 |
| comp | $Rd = (RSa == RSb); (RSa > RSb); (RSa < RSb)$ | 000001dddddaaaaaxxxxxssssss1100 |
| Immediate Instructions | | |
| addi | $Rd = RSa + Imm$ | 000011dddddaaaaaiiiiiiiiiii0000 |
| subi | $Rd = RSa - Imm$ | 000011dddddaaaaaiiiiiiiiiii0001 |
| addui | $Rd = RSa + Imm$ (Unsigned) | 000011dddddaaaaaiiiiiiiiiii0010 |
| subui | $Rd = RSa - Imm$ (Unsigned) | 000011dddddaaaaaiiiiiiiiiii0011 |
| andi | $Rd = RSa \& Imm$ | 000011dddddaaaaaiiiiiiiiiii0101 |
| ori | $Rd = RSa Imm$ | 000011dddddaaaaaiiiiiiiiiii0110 |
| xori | $Rd = RSa \oplus Imm$ | 000011dddddaaaaaiiiiiiiiiii0111 |
| sali | $Rd = RSa \ll Imm$ | 000011dddddaaaaaiiiiiiiiiii1000 |
| sari | $Rd = RSa \gg Imm$ | 000011dddddaaaaaiiiiiiiiiii1001 |
| slli | $Rd = RSa \ll Imm$ | 000011dddddaaaaaiiiiiiiiiii1010 |
| slri | $Rd = RSa \gg Imm$ | 000011dddddaaaaaiiiiiiiiiii1011 |
| compi | $Rd = (RSa == Imm); (RSa > Imm); (RSa < Imm)$ | 000011dddddaaaaaiiiiiiiiiii1100 |
| Load Instructions | | |
| lw | $Rd \leftarrow Imm(RSa)$ | 010010dddddaaaaa00iiiiiiiiiiii |
| lh | $Rd \leftarrow Imm(RSa)$ | 010010dddddaaaaa01iiiiiiiiiiii |
| lb | $Rd \leftarrow Imm(RSa)$ | 010010dddddaaaaa11iiiiiiiiiiii |
| lth | $Rd \leftarrow Imm(RSa)$ | 010010dddddaaaaa10iiiiiiiiiiii |

| Load Immediate Instructions | | |
|-----------------------------|---|-----------------------------------|
| li | (GPREGF) Rd = Imm | 000010dddd0000iiiiiiiiiiiiiiii |
| lsi | (SYSREGF) Rd = Imm | 000010dddd0001iiiiiiiiiiiiiiii |
| lgi | (GLREGF) Rd = Imm | 000010dddd0010iiiiiiiiiiiiiiii |
| li | (GPREGF) Rd = Imm | 000010dddd1000iiiiiiiiiiiiiiii |
| lsi | (SYSREGF) Rd = Imm | 000010dddd1001iiiiiiiiiiiiiiii |
| lgi | (GLREGF) Rd = Imm | 000010dddd1010iiiiiiiiiiiiiiii |
| Store Instructions | | |
| Instruction | Function | Binary |
| sw | RSb -> Imm(RSa) | 01101000iiiiiaaaabbbbbiiiiiiiiii |
| sh | RSb -> Imm(RSa) | 01101001iiiiiaaaabbbbbiiiiiiiiii |
| sb | RSb -> Imm(RSa) | 01101011iiiiiaaaabbbbbiiiiiiiiii |
| sth | RSb -> Imm(RSa) | 01101010iiiiiaaaabbbbbiiiiiiiiii |
| Jump Instructions | | |
| j | Next PC <- (R0 + Imm) | 000110iiiiiaaaaaiiiiiiiiiiiiiiii |
| jal | Next PC <- (R0 + Imm); RA0 <- PC | 000111iiiiiaaaaaiiiiiiiiiiiiiiii |
| jr | Next PC <- (RSa + Imm) | 000110iiiiiaaaaaiiiiiiiiiiiiiiii |
| jrl | Next PC <- (RSa + Imm); RA0 <- PC | 000111iiiiiaaaaaiiiiiiiiiiiiiiii |
| Branch Instructions | | |
| beq | Next PC <- (RSa == RSb) ? PC+Imm : PC+4 | 000101iiiiiaaaaabbbbbiiiiiiiiii00 |
| bne | Next PC <- (RSa != RSb) ? PC+Imm : PC+4 | 000101iiiiiaaaaabbbbbiiiiiiiiii01 |
| bgt | Next PC <- (RSa > RSb) ? PC+Imm : PC+4 | 000101iiiiiaaaaabbbbbiiiiiiiiii10 |
| blt | Next PC <- (RSa < RSb) ? PC+Imm : PC+4 | 000101iiiiiaaaaabbbbbiiiiiiiiii11 |
| System Instructions | | |
| syscall | System Call (Raise Privilege) | 00100dddddaaaaaa00000000iiiiiiii |
| sysret | System Return (Lower Privilege) | 00100dddddaaaaaa00000001iiiiiiii |
| stspr | (SYSRF) Rd <- RSa | 00100dddddaaaaaa00000010xxxxxxxx |
| ldspr | (SYSRF) RSa -> Rd | 00100dddddaaaaaa00000011xxxxxxxx |
| sync | Flush Pipeline | 00100xxxxxxxxxxx00000100xxxxxxxx |
| pmir | PML + Imm ? | 00100xxxxxxxxxxx00000111iiiiiiii |
| pmd | PML - Imm ? | 00100xxxxxxxxxxx00001000iiiiiiii |

6.4.1 Immediate

6.4.2 Load/Store

6.4.3 Branch/Jump

6.4.4 Co-Processor

7 Exceptions and Interrupts

7.1 Exceptions

7.2 Interrupts

7.2.1 User Level

7.2.2 Supervisor Level

Part IV

Co-Processors

8 Co-Processor Overview

Co-processors are the main point of the Fusion-Core architecture. As the ISA only defines the main core, the implementator is free to use whichever co-processors that would be necessary for an application. Hardware acceleration for vector instructions, encryption, floating point, communication, etc. There is no limitation for what kind of co-processor that could be used, only the number that could fit within the defined usable instructions.

It is important to note that the co-processors do not need to be of a RISC construction, due to this reliance on co-processors without specifying how they should be implemented. The main core is indeed RISC, as only the simplest instructions are defined that something as small as a microcontroller could use without issue.

The idea behind the separation of processor is to create different pipelines for each core. In doing so, the main core could be clocked faster than that of a pipeline requiring integer multiplication, or some other time consuming operation. As well, the individual cores could be clocked at their respective fastest frequencies, thus resulting in the fastest possible performance of each part. To deal with writing to memory with varying clock speeds for each core, a FIFO buffer is used to send each write to main memory. This FIFO is to create the illusion of write atomicity. The FIFO should have connections to allow for a seeming atomic read though, the programmer should note that reads require care; if there is a dependency on a slow core's written value. This programming paradigm is similar to that of parallel threads, where certain values may not be available until an unknown time. The ISA does not define any particular ways to handle this and it is either left up to the implementation, or programmer depending.

As shown in the instruction section, there are predefined regions for the co-processor slots. This is defined to allow for the compiler/assembler to work across implementations. At this time, only a fixed number of co-processors can be used on an implementation at a time, though future expansions to the ISA may change this.

8.1 Co-Processor Interface

The interface is designed to be extremely simple, as to make co-processors easy to implement. The inputs is just the output from the instruction fetch passthrough on the decode unit. As shown in the instruction list, there is a section of 6 bits for the co-processor's opcode. Please also note that this section of the instruction is different than the opcode, 0x3f (may change to 0x20 to allow for more opcodes for the co-processors), which indicates that the instruction is for the co-processor. To save connections, the normal opcode section is removed, leaving 26 bit physical instructions. Using more bits per instruction is not supported at this time, though it is possible to take the instruction fetch output directly.

8.1.1 Co-Processor Conventions

8.1.2 Register Connections

8.1.3 Decode unit Connections

8.2 Interface Connection Definitions

8.3 Adding custom Co-Processor

8.4 List of Co-Processors

8.4.1 Floating Point

8.4.2 System Unit

8.4.3 Memory Management Unit

8.4.4 Multiprocessor Communication Unit

9 Global Register File

The Global Register File is for simple message passing, and creating locks between co-processors. As it may be necessary to wait for a value to be computed by a co-processor, or lock specific parts of memory, the Global Register File creates an interface for ease of use between processing units.

This section is under development and will be updated to explain the connections and registers available.

10 Recommended Co-Processors

About This section will cover some basic coprocessors that have been approved and assigned coprocessor IDs. The full list of approved coprocessors will be included in a separate document.

10.1 Floating Point Unit

10.1.1 Registers

10.1.2 Instructions

10.2 System Unit

10.2.1 Registers

10.2.2 Instructions

10.3 Memory Management Unit

10.3.1 Registers

10.3.2 Instructions

10.4 Multiprocessor Communication Unit

10.4.1 Registers

10.4.2 Instructions